

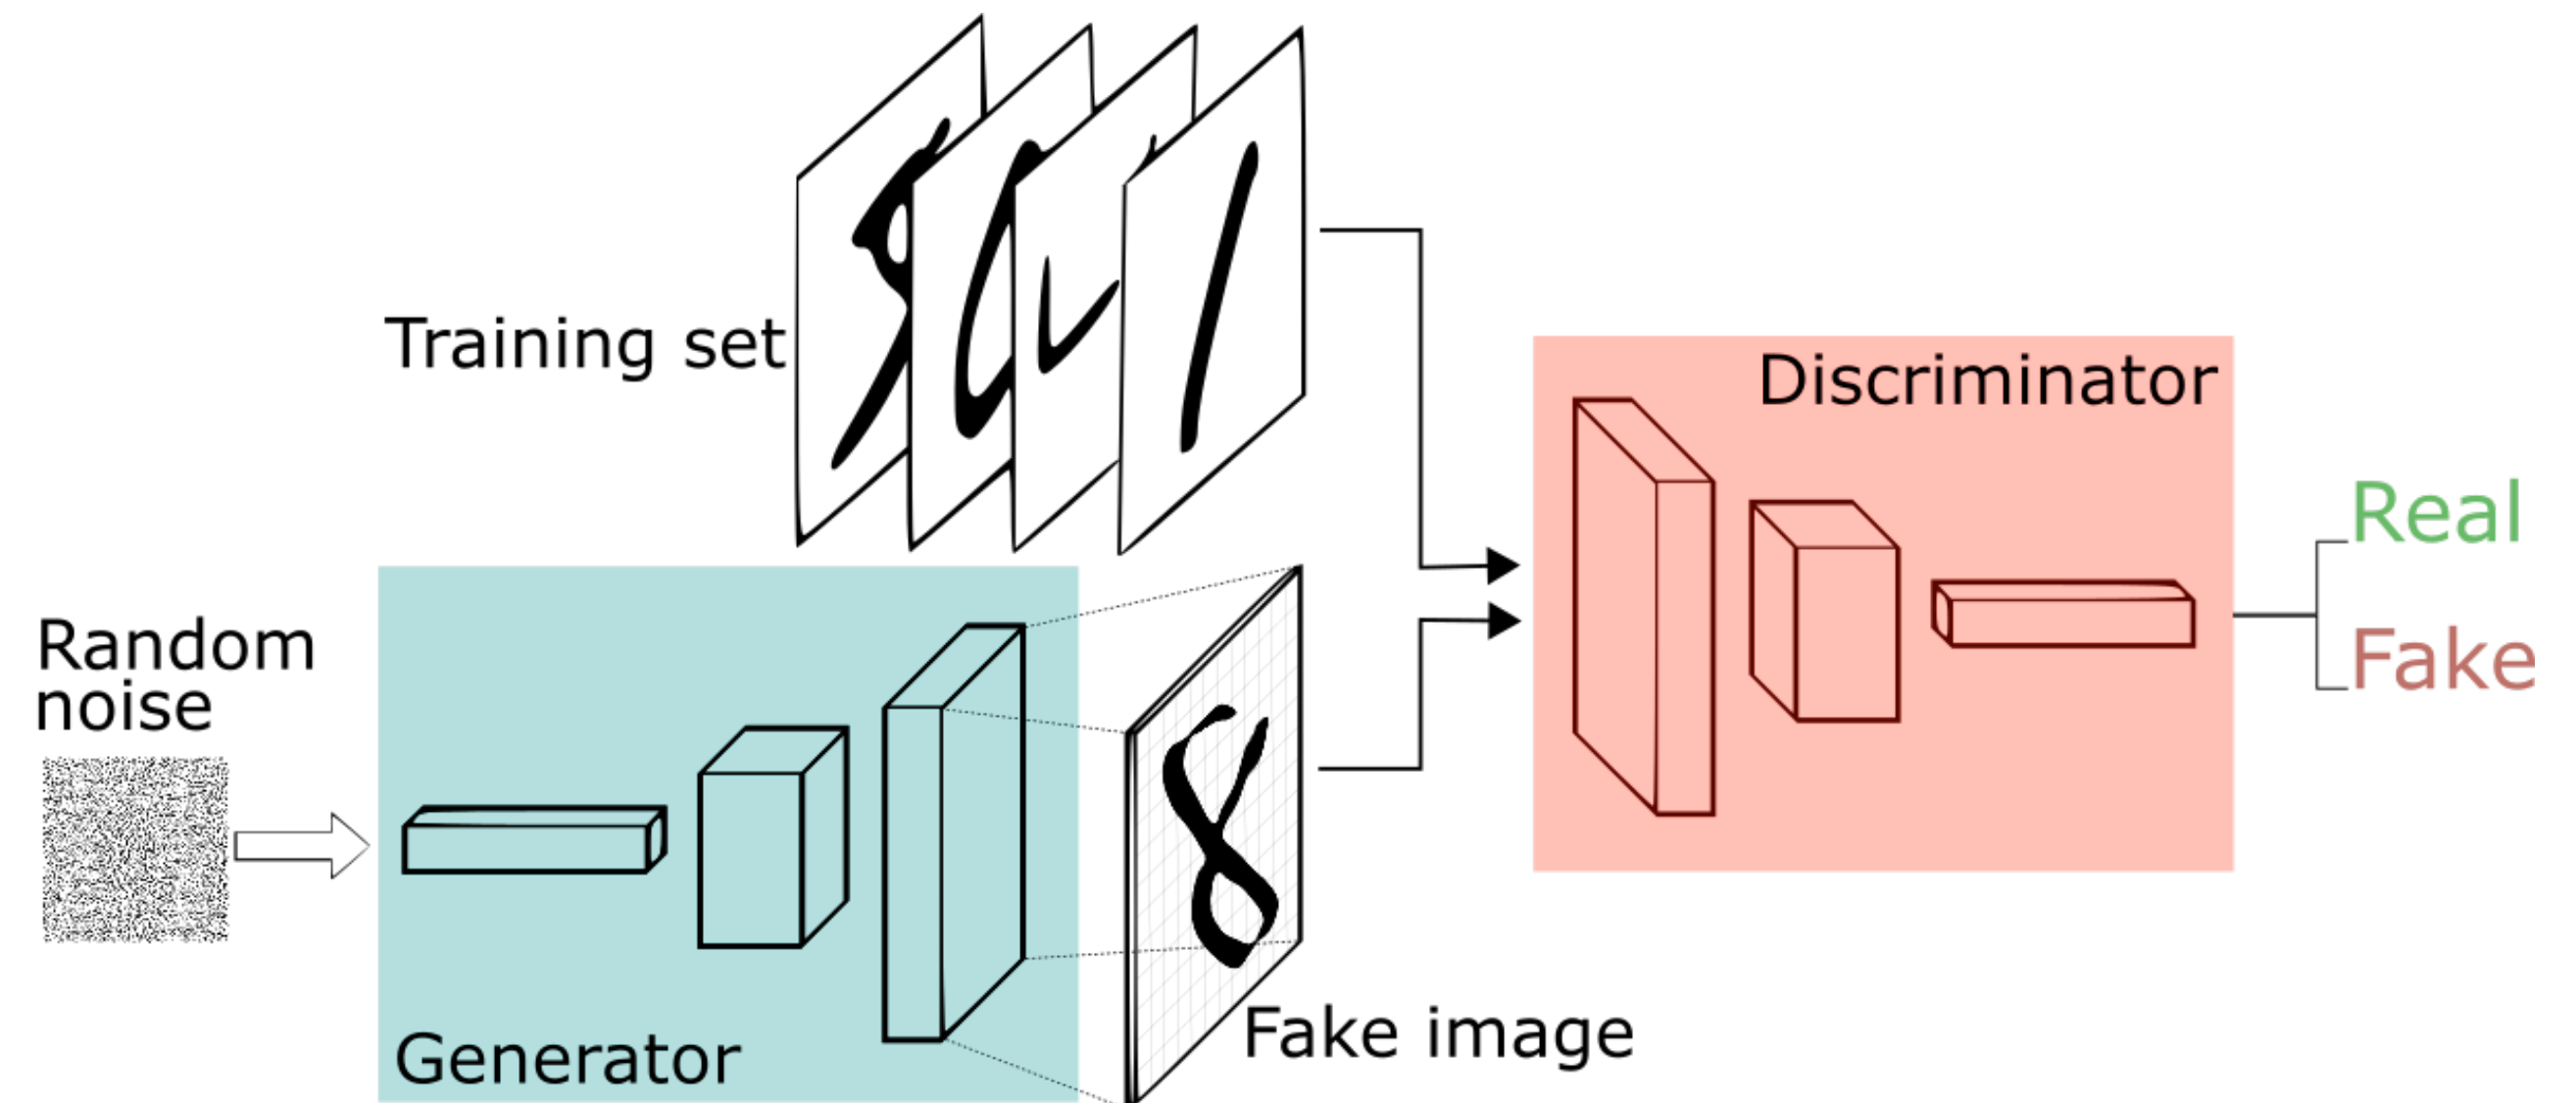
# Generative Adversarial Networks (GANs)

# Analogy

- Generative: team of counterfeiters, trying to fool police with counterfeit money
- Discriminative: police trying to detect the counterfeit money
- Competition drives to improve both, until counterfeits are virtually indistinguishable from genuine currency.
- Now counterfeiters have, as a side effect, learned something about real currency.

# Big Idea

- Train a generative model  $G(z)$  to generate data with random noise,  $z$ , as input
- Adversary is discriminator  $D(x)$  which is trained to distinguish synthetic and true data
- Represent  $G(z)$  &  $D(z)$  as multilayer perceptrons for differentiability





```
def make_generator_model():
    model = tf.keras.Sequential()
    model.add(layers.Dense(7*7*256, use_bias=False, input_shape=(100,)))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Reshape((7, 7, 256)))
    assert model.output_shape == (None, 7, 7, 256) # Note: None is the batch size

    model.add(layers.Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same', use_bias=False))
    assert model.output_shape == (None, 7, 7, 128)
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same', use_bias=False))
    assert model.output_shape == (None, 14, 14, 64)
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same', use_bias=False,
                                     activation='tanh'))
    assert model.output_shape == (None, 28, 28, 1)

    return model
```



```
def make_discriminator_model():  
    model = tf.keras.Sequential()  
    model.add(layers.Conv2D(64, (5, 5), strides=(2, 2), padding='same',  
                             input_shape=[28, 28, 1]))  
  
    model.add(layers.LeakyReLU())  
    model.add(layers.Dropout(0.3))  
  
    model.add(layers.Conv2D(128, (5, 5), strides=(2, 2), padding='same'))  
    model.add(layers.LeakyReLU())  
    model.add(layers.Dropout(0.3))  
  
    model.add(layers.Flatten())  
    model.add(layers.Dense(1))  
  
    return model
```

$$\min_G \max_D E_{z,x} \left[ \log D(G(z)) + \log(1 - D(x)) \right]$$

*Probability that discriminator identifies real data as real*  

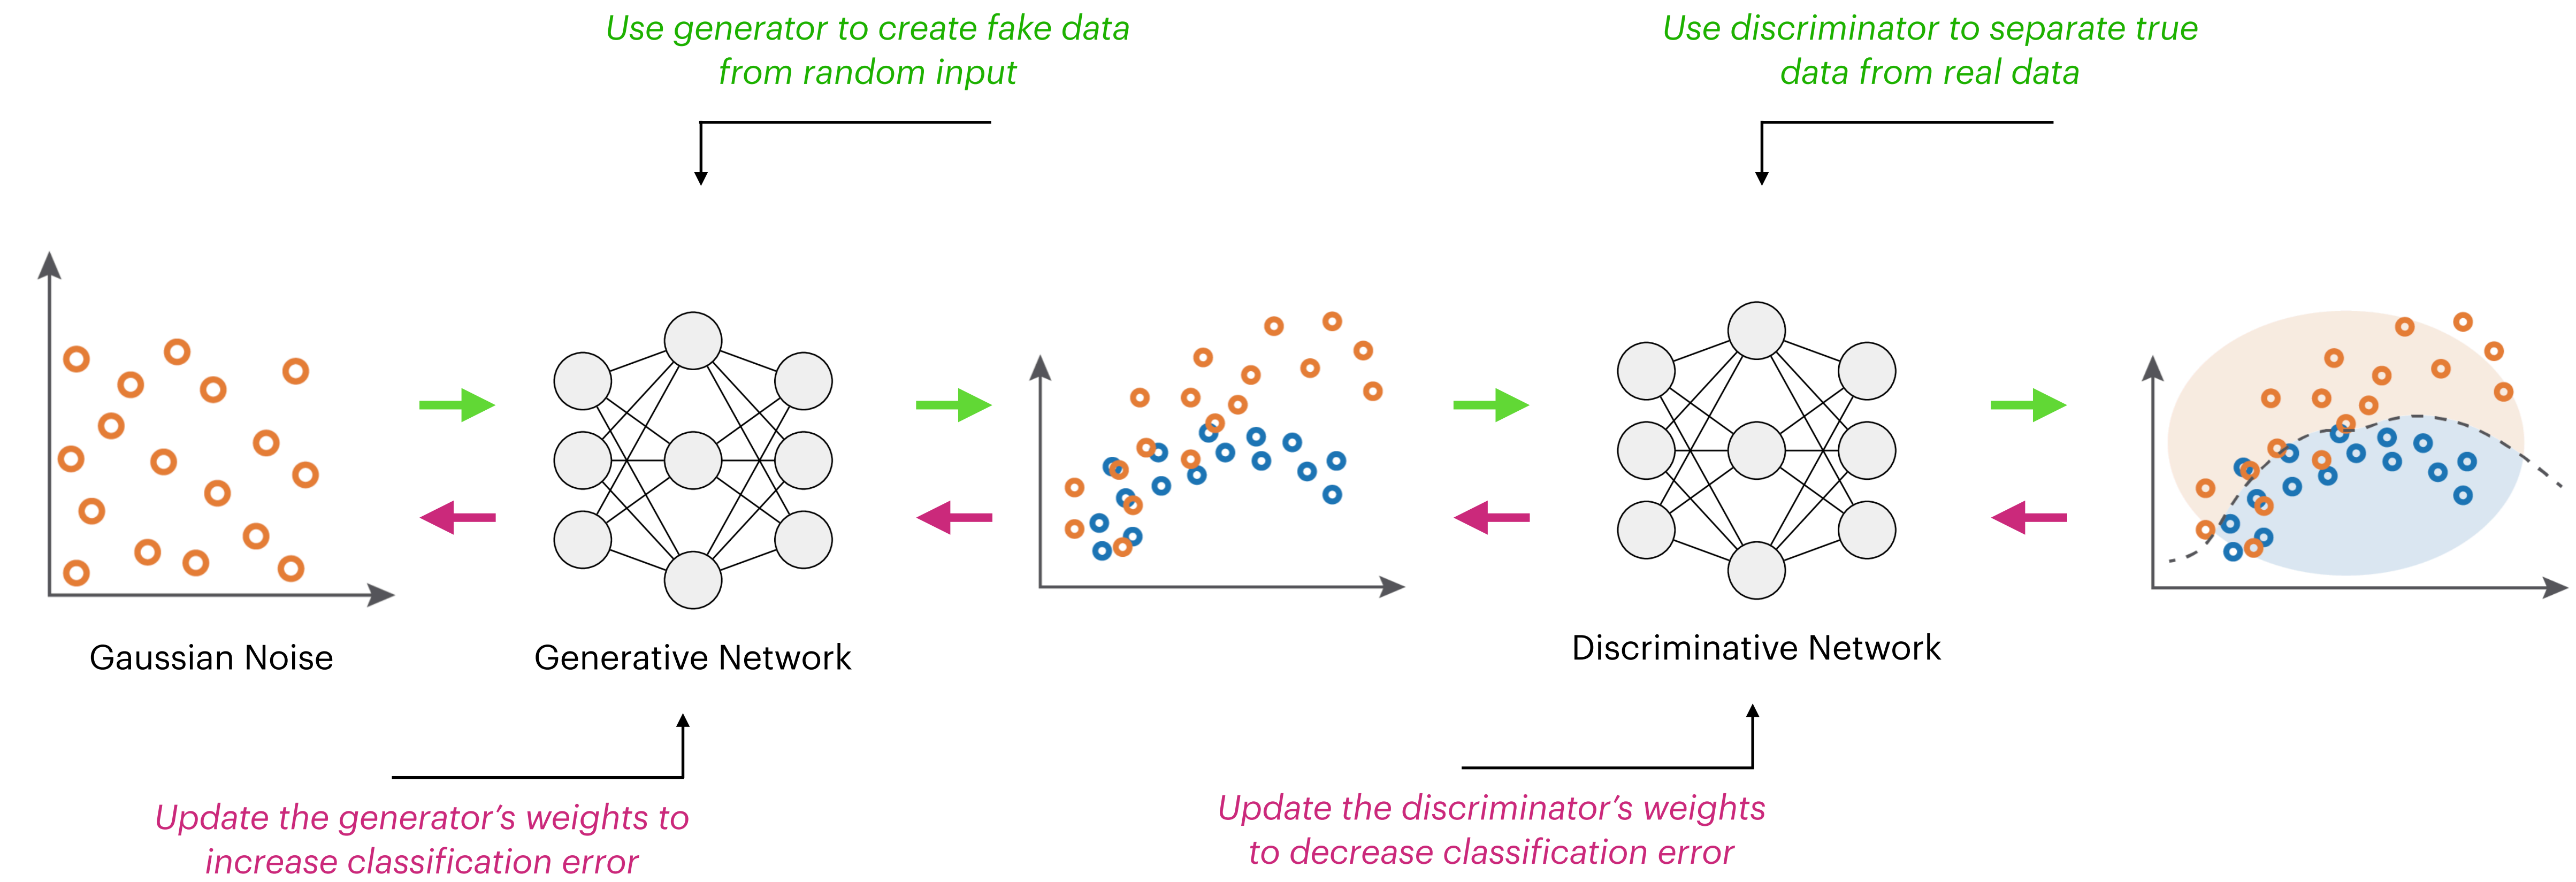

---

*Probability that discriminator identifies generated image as fake*

**Global Optimum: A global minimum of this minimax function**

# Training GANs





Input random variable  
(Drawn from a simple  
distribution, e.g gaussian  
noise)

The generative network is  
trained to **maximise** the  
final classification error

The **generated  
distribution** and **true  
distribution** are not  
compared directly

The discriminative network  
is trained to **minimise** the  
final classification error

The classification error is  
the reference metric for  
the training of both  
networks




# Drawbacks of using GANs

# Mode Collapse

- Natural data distributions are highly complex and multimodal.
- During mode collapse, the generator produces samples that belong to a limited set of modes. The generator believes that it can fool the discriminator by locking on to a single mode, so it produces only samples from this exclusive mode.
- The discriminator eventually figures this out. Generator just locks on to another mode.
- Cycle repeats

Without mode collapse

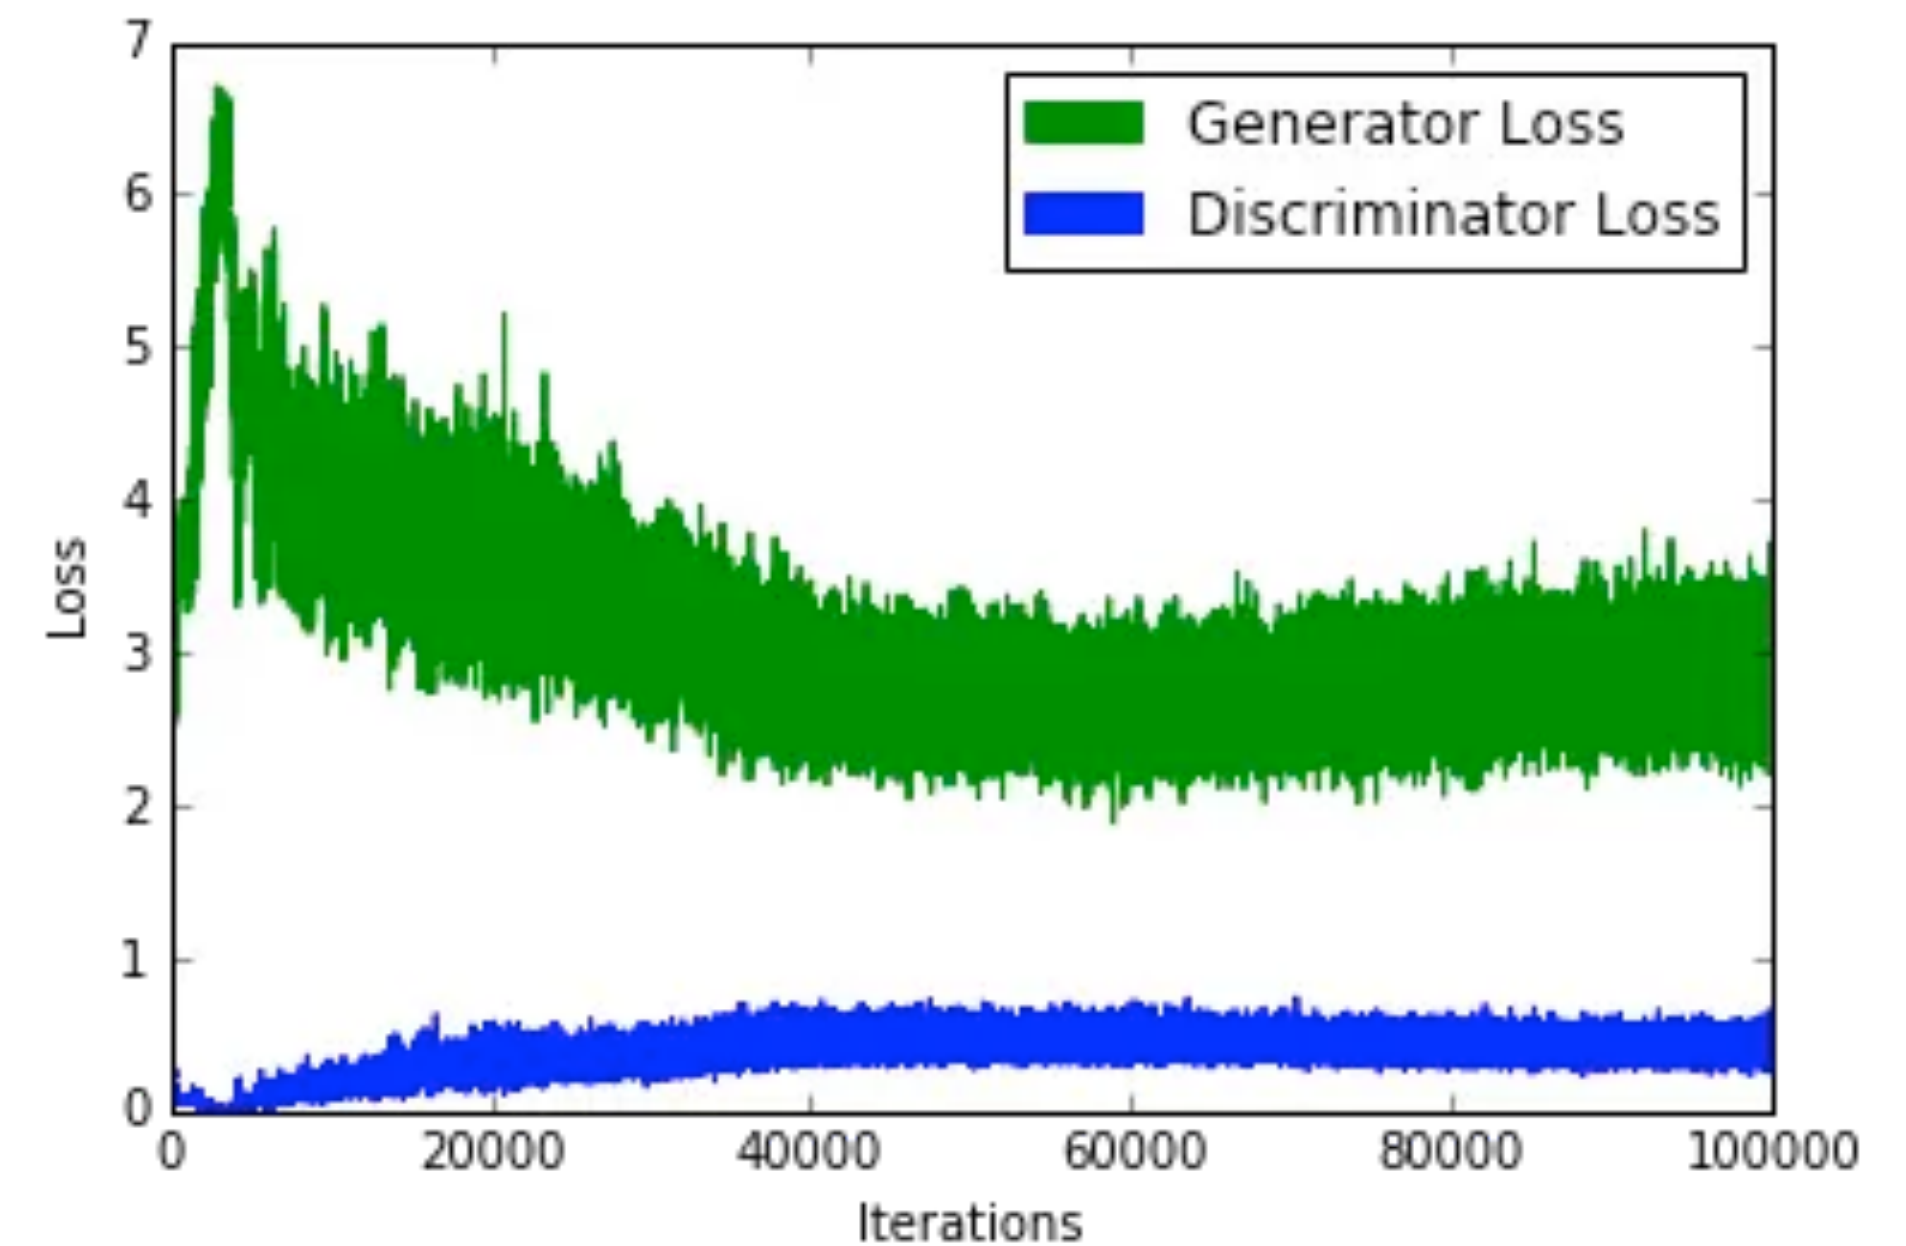


With mode collapse

# Convergence

## When do we stop training?

Since the generator loss improves when the discriminator loss degrades (and vice-versa), we cannot judge convergence based on the value of the loss function.



Typical plot of a GAN loss function. Note how convergence cannot be interpreted from this plot

# Quality

- It is difficult to quantitatively tell when the generator is producing high quality samples.
- Additional perceptual regularisation added to the loss function can help mitigate the situation to some extent.

# Metrics

- The GAN objective function explains how well the Generator or the Discriminator is performing with respect to its adversary.
- It does not represent the quality or diversity of its output.
- Hence we need distinct metrics that can measure these.

# Techniques for improving Performance

- Alternative loss functions: Replace Jensen Shannon divergence of conventional GANs with Wasserstein Loss
- Two Timescale Update Rule (TTUR): Use a different learning rate for D and G
- Gradient penalty: greatly enhances stability and reduces mode collapse
- Stacking GANs: Use multiple GANs to solve an easier version of the problem

# Recent advances in GANs


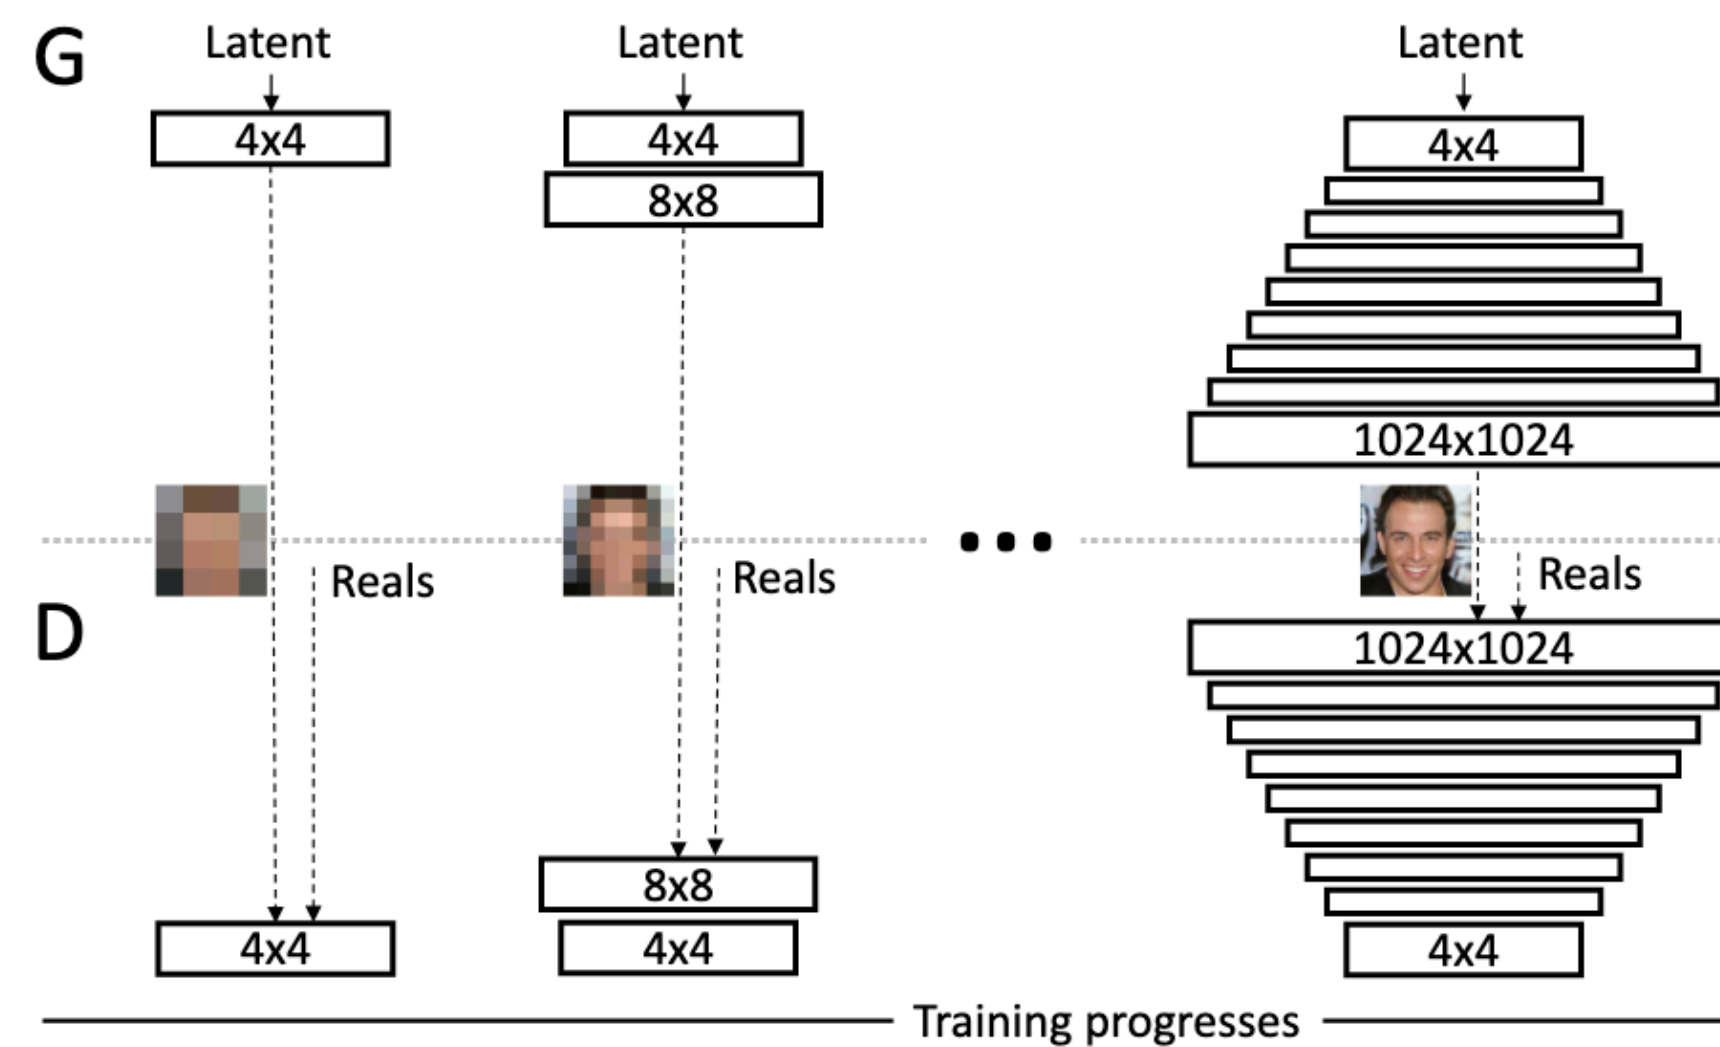


# Progressive growing of GANs

A type of GAN that involves the conditional generation of images by a generator model

## Advantages:

1. Produce images of unprecedented quality
2. Speeds up training, and also stabilises it.




# Conditional GANs

A type of GAN that involves the conditional generation of images by a generator model

## Advantages:

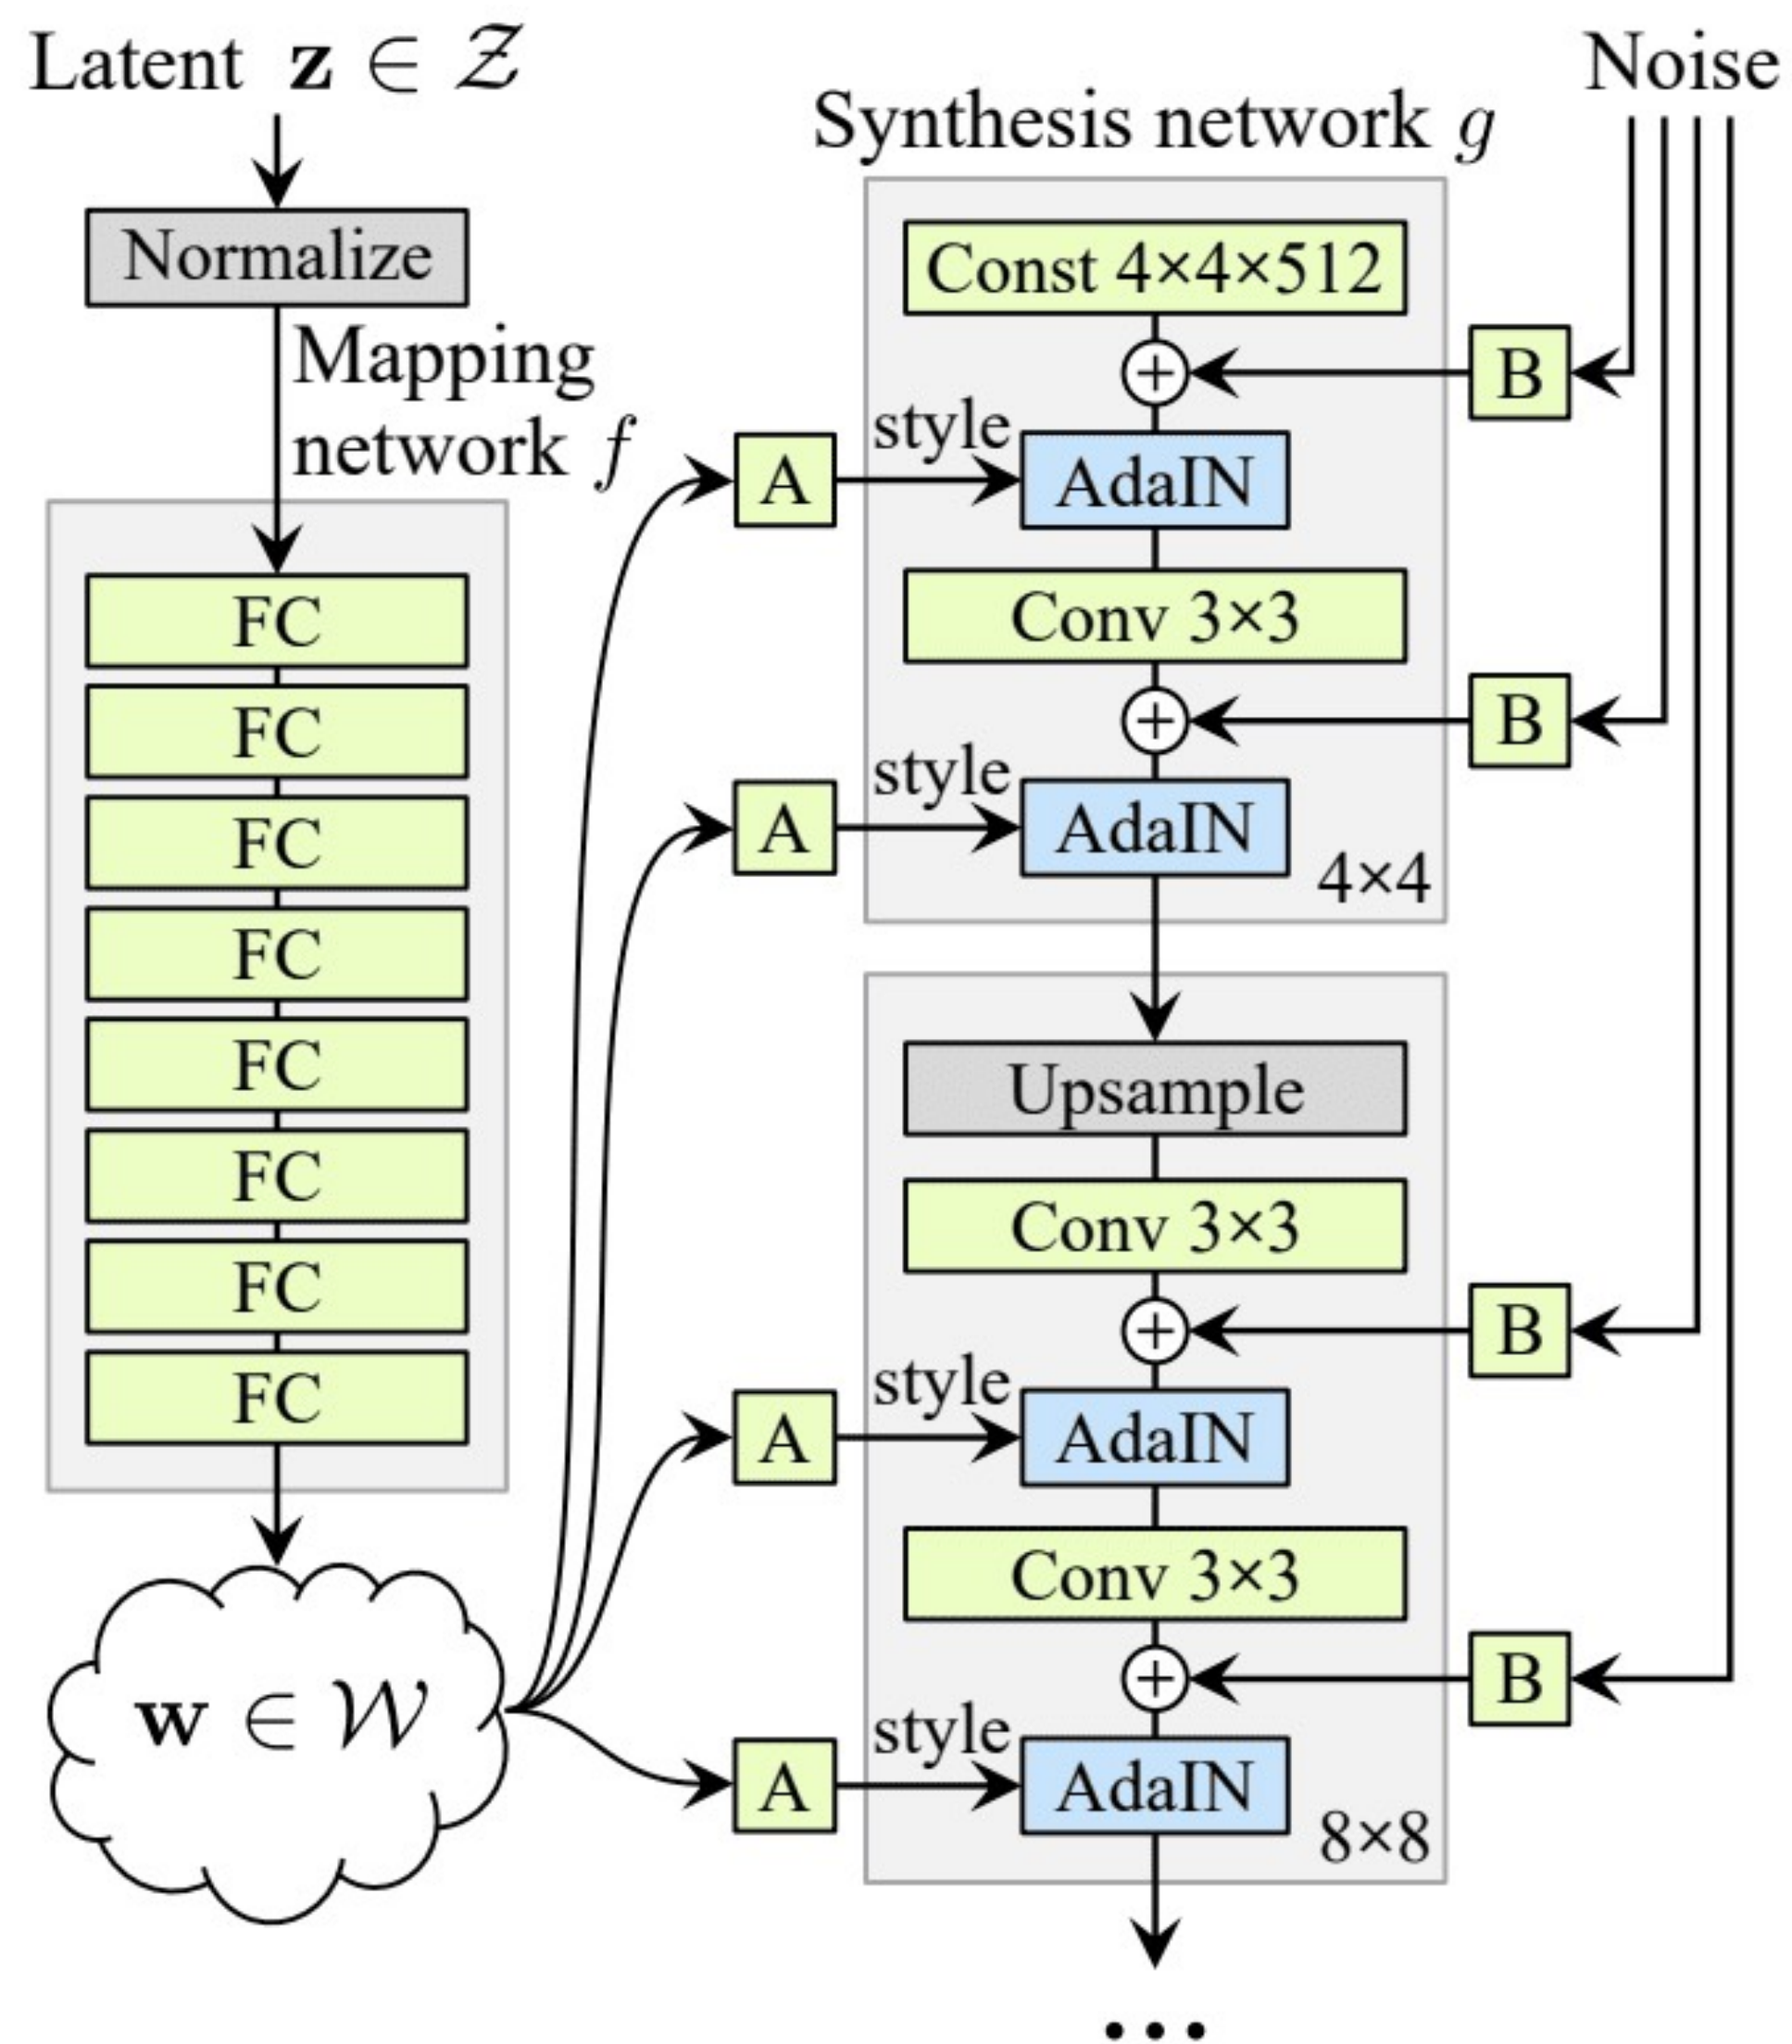
1. Convergence will be faster. Even the random distribution that the fake images will follow will have some pattern
2. You can control the output of the generator at test time, by giving the label for the image you want to generate.





# StyleGAN

Progressive growing + Style transfer

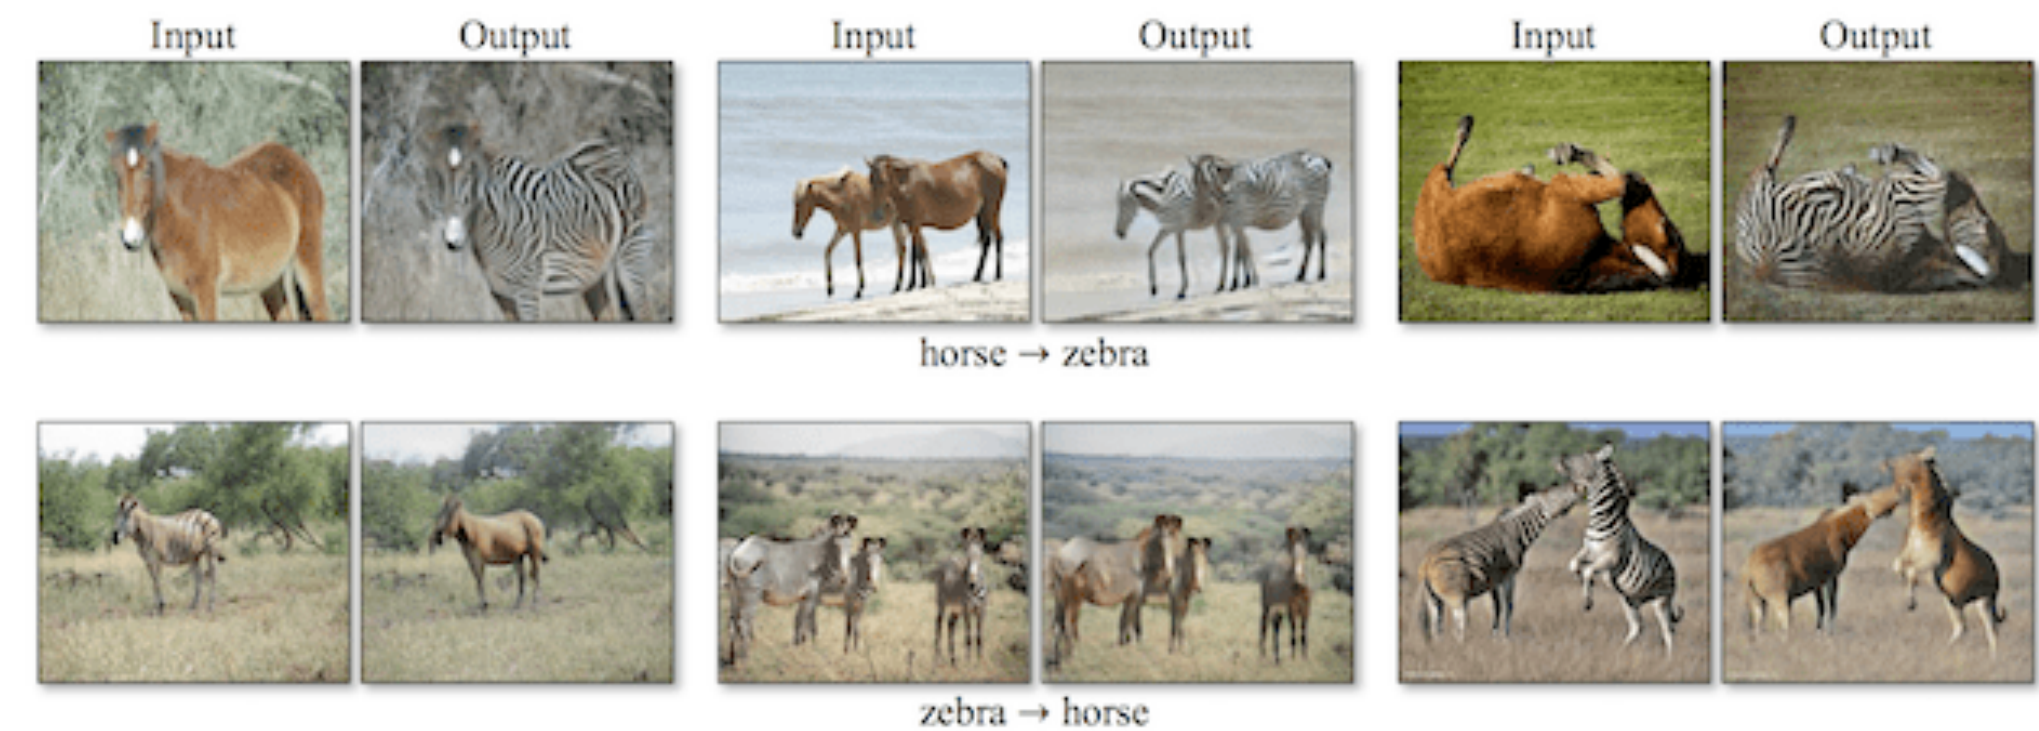




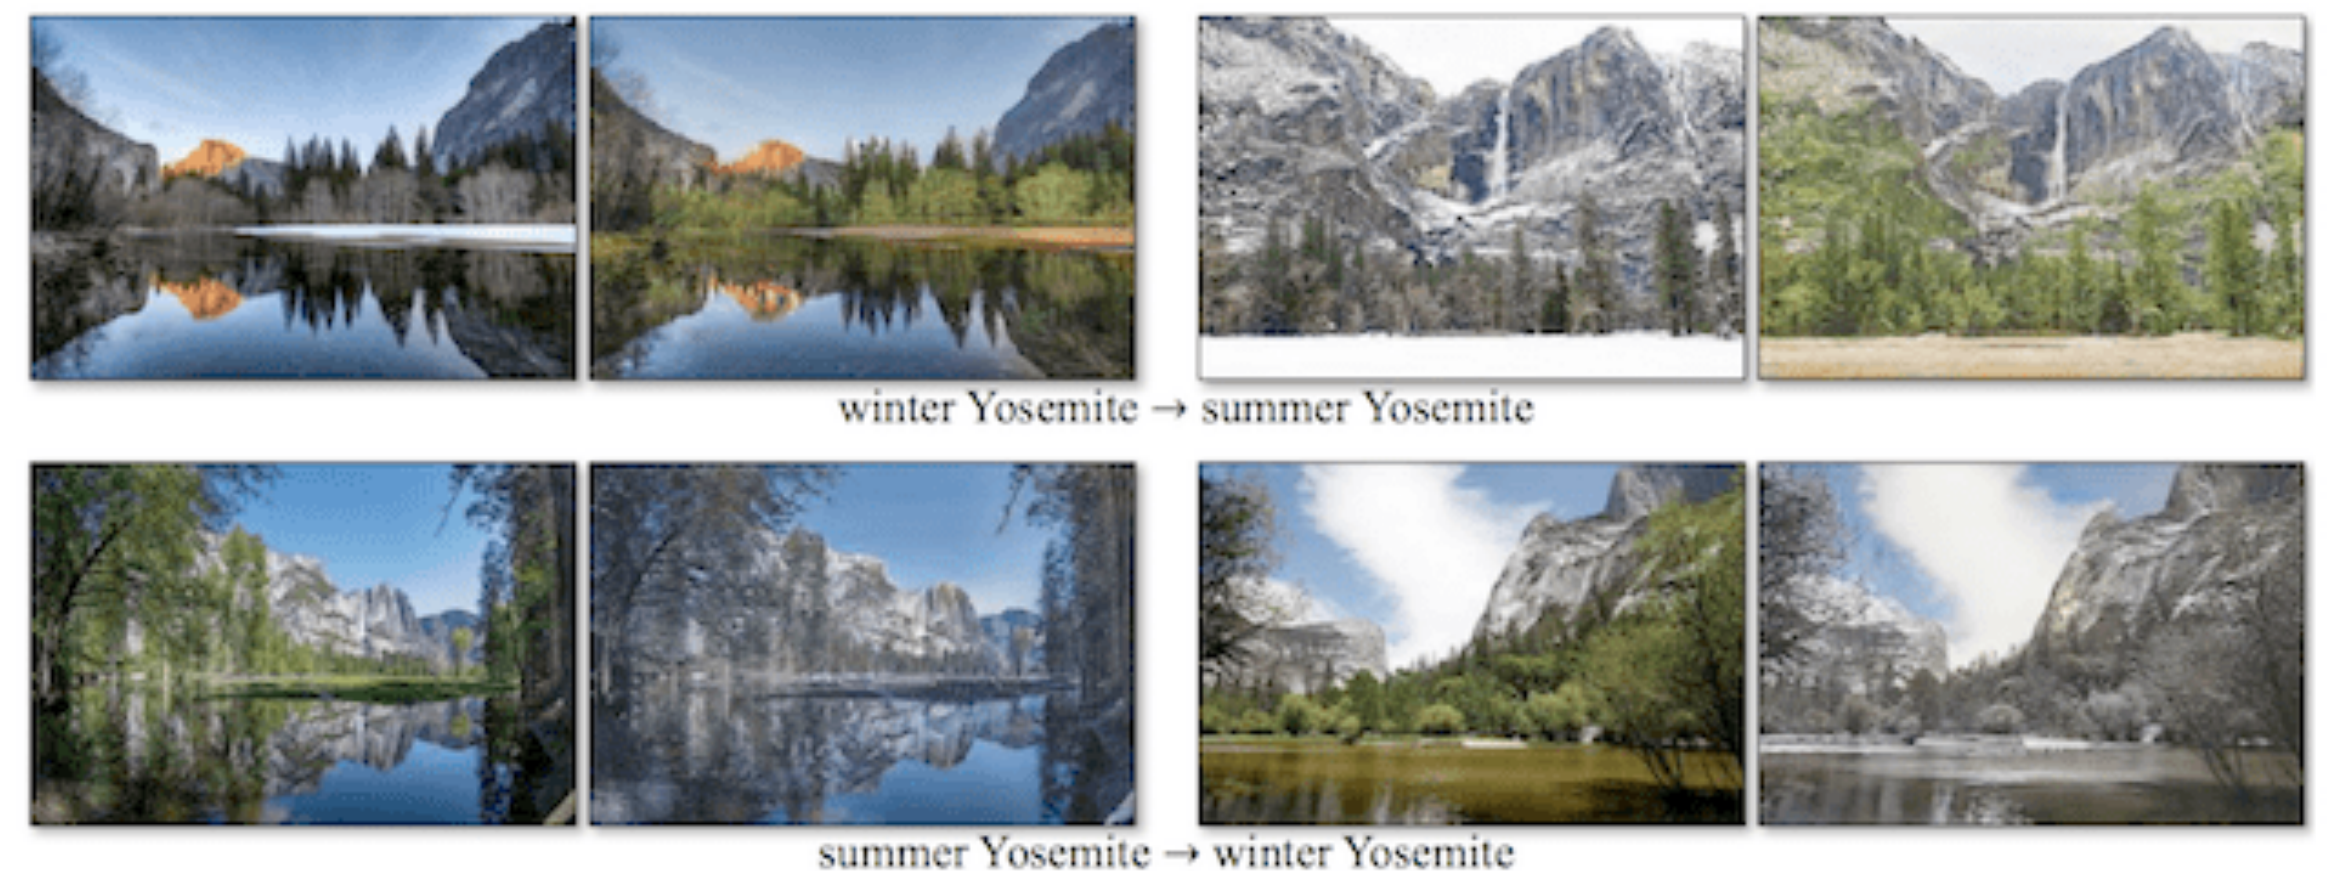
# CycleGAN

## For image-to-image translation without paired examples

- Image-to-image translation often involves the controlled modification of an image and requires large datasets of paired images.
- CycleGAN is a technique for training unsupervised image translation models via GAN architecture using unpaired collections of images from different domains
- Can be used for season translation, object transfiguration, style transfer, photograph enhancement.



Object transfiguration



Season transfer

# References

- [ARTICLE]: Understanding Generative Adversarial Networks
- [SLIDES]: Foundations of GANs
- [ARTICLE]: Advances in GANs
- [ARTICLE]: Introduction to StyleGANs
- [PAPER]: Progressive growing of GANs
- [ARTICLE]: CycleGAN for image translation